Diego Villagran Salazar

Machine Learning Engineer | MLOps | Applied AI

diegovillasal@gmail.com | +52 55 7207 1183 | www.dvillagrans.dev | LinkedIn | GitHub

Professional Summary

Data Science student with hands-on experience building machine learning pipelines, model training and serving, and data processing APIs in production. Skilled in Python (Pandas, NumPy, Scikit-learn, TensorFlow), feature engineering, evaluation, and containerized deployment with FastAPI/Flask and Docker. Delivered 85%+ accuracy on EEG-based emotion classification with real-time streaming under 100ms, and supported 10K+ daily requests across services. Seeking an ML Engineer role to scale training, inference, and monitoring in production.

Education

Bachelor of Science in Data Science

Superior School of Computer Science (ESCOM-IPN) — *2022 – 2026 (December)*

Mexico City, Mexico

- Maintained GPA of 8.5/10.0 while completing 45+ credit hours in advanced coursework
- Specialized in Machine Learning, Statistical Modeling, Data Engineering, and Big Data Analytics

Professional Experience

AI & Automation Intern

EyeNet — *April 2025 – Present*

Remote

- Developed LLM-assisted document extraction pipelines integrating OpenAI and Gemini APIs with Python, structuring 200+ documents per day with 92% extraction accuracy
- Built data processing APIs and services generating model-ready datasets on PostgreSQL, Redis, and MongoDB, supporting 10K+ daily requests
- Containerized data and model services with Docker and implemented CI/CD and telemetry, reducing deployment time by 50% and detecting 95%+ issues pre-production
- Created evaluation and monitoring dashboards with Grafana and Prometheus to track throughput, error rates, and data quality

Projects & Research

Qalma - Applied ML on EEG Signals

Highlighted Project — 2025

- Built end-to-end ML pipeline for EEG signals at 256 Hz: preprocessing, feature extraction, and classification achieving 85%+ accuracy
- Implemented real-time streaming and inference with latency under 100ms; stored datasets and predictions in Supabase/PostgreSQL
- Served models via FastAPI/Flask with Docker; added batch jobs for historical analysis and drift checks

Microservices Orchestration & Monitoring

MLOps & Platform Project — 2025

- Deployed 8+ services (APIs, PostgreSQL, Redis, Nginx) with Docker Compose, health checks, and automated recovery achieving 99.5% uptime
- Built observability with Grafana/Prometheus exposing 25+ metrics; performed load testing (Apache Bench) reducing latency by 40%

Technical Skills

Machine Learning: Python (Pandas, NumPy), Scikit-learn, TensorFlow, Feature Engineering, Model Training/Evaluation, Time-series Processing
MLOps & Serving: FastAPI/Flask, Docker, CI/CD, Monitoring (Grafana, Prometheus), Model Packaging
Data & SQL: SQL, PostgreSQL, MongoDB, Redis, ETL Pipelines, Data Modeling
Cloud & Tools: Git, AWS, Azure, Linux Administration, Shell Scripting
Analytics: A/B Testing, Statistical Modeling, Power BI, Tableau

Certifications & Languages

Certifications: Google AI Essentials (Coursera), Build a Data Warehouse with BigQuery, Engineer Data for Predictive Modeling with BigQuery ML, Manage Kubernetes in Google Cloud
Languages: Spanish (Native proficiency), English (B2 in-course)